



# Code base Seminar

샘이나 샘이나~

Backend.ai 멘티 이유찬

# Order

Step 1

Overview

Step 2

Client

Step 3

Manager

Step 4

Agent



01

OVERVIEW

# Backend.ai Tutorial

```
backend.ai run -c 'print("hello")' python
```

```
uchanlee — ubuntu@CA4: ~/uchan/backend.ai/backend.ai-dev — ssh - sh enter.sh — 109x36

ream [2878292] STREAM_EXECUTE(ak:AKIAIOSFODNN7EXAMPLE,
s:91100b1b-8301-468d-bd06-2f51c8a9599f)
2021-10-16 14:42:06.192 INFO aiohttp.access [2878292]
127.0.0.1 [16/Oct/2021:05:42:06 +0000] "GET /stream/se
ssion/91100b1b-8301-468d-bd06-2f51c8a9599f/execute HTT
P/1.1" 101 0 "-" "Backend.AI Client for Python 21.09.0
a2"

2021-10-16 14:42:06.181 INFO ai.backend.agent.server [
2886242] rpc::execute(k:91100b1b-8301-468d-bd06-2f51c8
a9599f, run-id:5b23b4c60834dbb0, mode:query, code:'pri
nt("hello")')
2021-10-16 14:42:06.190 DEBUG ai.backend.agent.agent [
2886242] _execute(91100b1b-8301-468d-bd06-2f51c8a9599f
) finished

1] started.
2021-10-16 14:41:10 INFO ai.backend.web.server [288975
3] started.
2021-10-16 14:41:10 INFO ai.backend.web.server [288975
2] started.
2021-10-16 14:41:10 INFO ai.backend.web.server [288975
4] started.

INFO Started service.
INFO Backend.AI Storage Proxy
INFO Runtime: Python 3.9.5 (env: /home/ubuntu/.pyenv/v
ersions/venv-4d9o0zwu-storage-proxy)
INFO Node ID: i-storage-proxy-01
DEBUG debug mode enabled.
INFO Using uvloop as the event loop backend

(venv-4d9o0zwu-client) ubuntu@CA4:~/uchan/backend.ai/backend.ai-dev/client-py$ backend.ai run -c 'print("hell
o")' python
- Session name prefix: pysdk-e197afd41d
✓ [0] Session pysdk-e197afd41d is ready (domain=default, group=default).
hello
✓ [0] Execution finished. (exit code = 0)
(venv-4d9o0zwu-client) ubuntu@CA4:~/uchan/backend.ai/backend.ai-dev/client-py$ █

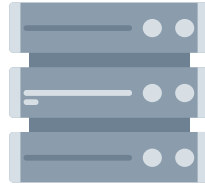
[11] 0:user* "CA4" 14:41 16-Oct-21
```

# Components

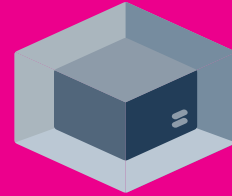
Client



Manager



Agent





02

Client

# Client CLI

```
def _create_cmd(docs: str = None):

    @click.argument("image")
    @click.option('-t', '--name', '--client-token', metavar='NAME',
                  help='Specify a human-readable session name. '
                        'If not set, a random hex string is used.')
    @click.option('-o', '--owner', '--owner-access-key',
                  metavar='ACCESS_KEY',
                  help='Set the owner of the target session explicitly.')

    # job scheduling options
    @click.option('--type', metavar='SESSTYPE',
                  type=click.Choice(['batch', 'interactive']),
                  default='interactive',
                  help='Either batch or interactive')

    @click.option('--starts-at', metavar='STARTS_AT', type=str,
                  default=None,
                  help='Let session to be started at a specific or relative
time.')
    ...
```

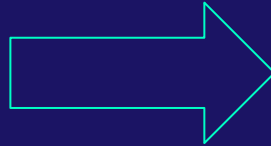
```
def create(
    # base args
    image: str,
    name: str | None,
    owner: str | None,
    ...
    ...
) -> None:

    ...

    with Session() as session:
        try:
            compute_session = session.ComputeSession.get_or_create(
                image,
                name=name,
                type_=type,
                starts_at=starts_at,
                enqueue_only=enqueue_only,
                max_wait=max_wait,
                no_reuse=no_reuse,
                ...
            )
```

# Client Func(API)

```
async def get_or_create(
    cls,
    image: str = *,
    name: str = None,
    type_: str = 'interactive',
    starts_at: str = None,
    enqueue_only: bool = False,
    max_wait: int = 0,
    no_reuse: bool = False,
    mounts: List[str] = None,
    mount_map: Mapping[str, str] = None,
    envs: Mapping[str, str] = None,
    startup_command: str = None,
    resources: Mapping[str, int] = None,
    resource_opts: Mapping[str, int] = None,
    cluster_size: int = 1,
    cluster_mode: Literal['single-node', 'multi-node'] = 'single-node',
    domain_name: str = None,
    ...,
) -> ComputeSession:
```



1. Set Param
2. Request & response

```
rqst = Request('POST', f!/{prefix}')

params: Dict[str, Any] = {
    'tag': tag,
    get_naming(api_session.get().api_version, 'name_arg'): name,
    'config': {
        ...
    },
}

if api_session.get().api_version >= (6, '20200815'):
    params['clusterSize'] = cluster_size
    params['clusterMode'] = cluster_mode
...

rqst.set_json(params)
async with rqst.fetch() as resp:
    data = await resp.json()

    o = cls(name, owner_access_key) # type: ignore
    if api_session.get().api_version[0] >= 5:
        o.id = UUID(data['sessionId'])
    ...

    return o
```



```

    async def get_or_create(
        cls,
        ...
        ...
    ) -> ComputeSession:
        ...

        rqst = Request('POST', f/{prefix}')
        params: Dict[str, Any] = {
            'tag': tag,
            'get_naming(api_session.get().api_version, 'name_arg')': name,
            'config': {
                ...
            },
        }
        if api_session.get().api_version >= (6, '20200815'):
            params['clusterSize'] = cluster_size
            params['clusterMode'] = cluster_mode
            ...

        rqst.set_json(params)
        async with rqst.fetch() as resp:
            data = await resp.json()
            o = cls(name, owner_access_key) # type: ignore
            if api_session.get().api_version[0] >= 5:
                o.id = UUID(data['sessionId'])
            ...
        return o

```



# Manager

```

    async def get_or_create(
        cls,
        ...
        ...
    ) -> ComputeSession:
        ...

        rqst = Request('POST', f/{prefix}')
        params: Dict[str, Any] = {
            'tag': tag,
            'get_naming(api_session.get().api_version, 'name_arg')': name,
            'config': {
                ...
            },
        }
        if api_session.get().api_version >= (6, '20200815'):
            params['clusterSize'] = cluster_size
            params['clusterMode'] = cluster_mode
        ...

        rqst.set_json(params)
        async with rqst.fetch() as resp:
            data = await resp.json()
            o = cls(name, owner_access_key) # type: ignore
            if api_session.get().api_version[0] >= 5:
                o.id = UUID(data['sessionId'])
            ...
        return o

```



03

Manager

# Manager API session

```
async def _create(request: web.Request, params: Any) ->
web.Response:
    if params['domain'] is None:
        params['domain'] = request['user']['domain_name']
    scopes_param = {
        'owner_access_key': (None if params['owner_access_key'] is
undefined
                             else params['owner_access_key'])
    }
    requester_access_key, owner_access_key = await
get_access_key_scopes(request, scopes_param)
    log.info('GET_OR_CREATE (ak:{0}/{1}, img:{2}, s:{3})',
            requester_access_key, owner_access_key if
owner_access_key != requester_access_key else '**',
            params['image'], params['session_name'])

    root_ctx: RootContext = request.app['_root.context']
    ...
```

```
...
# Resolve the image reference.
try:
    requested_image_ref = \
        await ImageRef.resolve_alias(params['image'],
root_ctx.shared_config.etcd)
    async with root_ctx.db.begin() as conn:
        query = (sa.select([domains.c.allowed_docker_registries])
                .select_from(domains)
                .where(domains.c.name == params['domain']))
        allowed_registries = await conn.scalar(query)
        if requested_image_ref.registry not in allowed_registries:
            raise AliasResolutionFailed
except AliasResolutionFailed:
    raise ImageNotFound('unknown alias or disallowed registry')
...
```

# Manager API session

```
...
# Resolve the image reference.
try:
    requested_image_ref = \
        await ImageRef.resolve_alias(params['image'],
root_ctx.shared_config.etcdd)
    async with root_ctx.db.begin() as conn:
        query = (sa.select([domains.c.allowed_docker_registries])
            .select_from(domains)
            .where(domains.c.name == params['domain']))
        allowed_registries = await conn.scalar(query)
        if requested_image_ref.registry not in allowed_registries:
            raise AliasResolutionFailed
except AliasResolutionFailed:
    raise ImageNotFound('unknown alias or disallowed registry')
...
```

```
try:

    kernel_id = await
asyncio.shield(root_ctx.registry.enqueue_session(
    session_creation_id,
    params['session_name'], owner_access_key,
    [{
        'image_ref': requested_image_ref,
        ...
    }],
    params['config']['scaling_group'],
    params['session_type'],

    ...
))
resp['sessionId'] = str(kernel_id) # changed since API v5
resp['sessionName'] = str(params['session_name'])
resp['status'] = 'PENDING'
resp['servicePorts'] = []
resp['created'] = True
```

# Manager registry

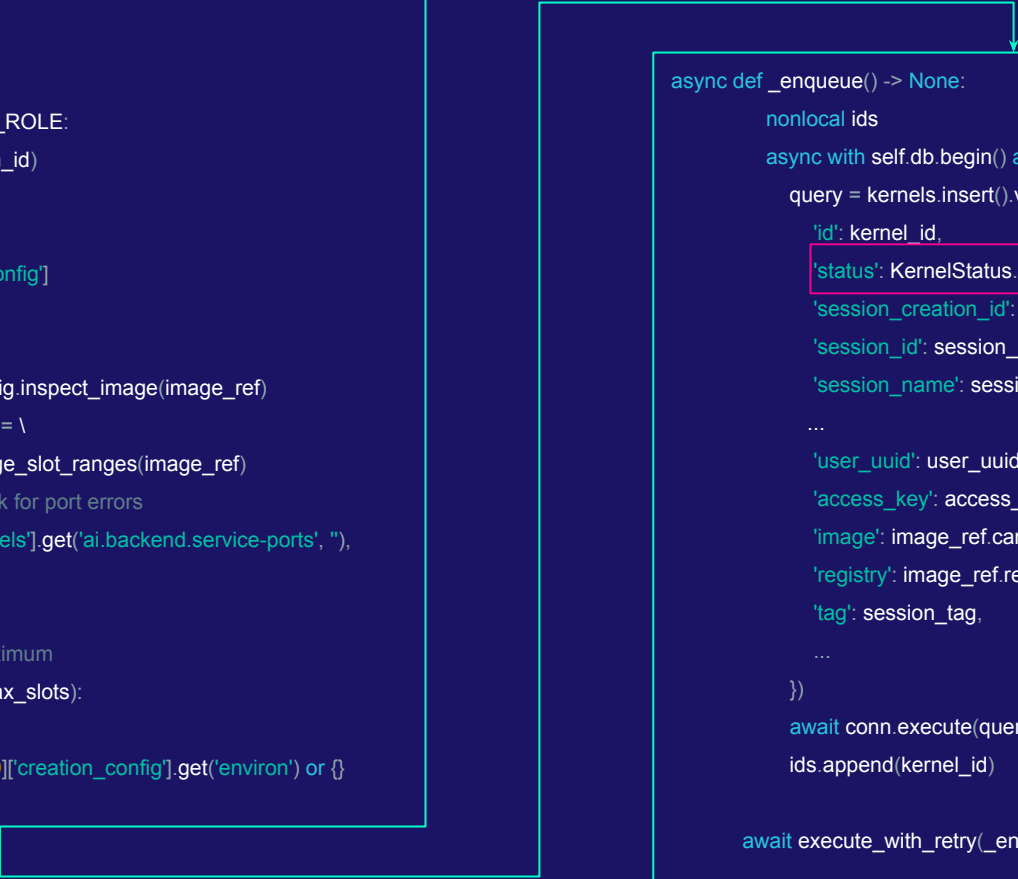
```
async def enqueue_session(
    self,
    session_creation_id: str,
    session_name: str,
    ...
) -> SessionId:

    mounts =
kernel_enqueue_configs[0]['creation_config'].get('mounts') or []
    mount_map =
kernel_enqueue_configs[0]['creation_config'].get('mount_map') or {}
    session_id = SessionId(uuid.uuid4())
    ...
```

```
...
is_multicontainer = cluster_size > 1
if is_multicontainer:
    if len(kernel_enqueue_configs) == 1:
        log.debug(
            'enqueue_session(): replicating kernel_enqueue_config with
cluster_size={},
            cluster_size,
        )
        # the first kernel_config is repliated to sub-containers
        assert kernel_enqueue_configs[0]['cluster_role'] ==
DEFAULT_ROLE
        kernel_enqueue_configs[0]['cluster_idx'] = 1
        for i in range(cluster_size - 1):
            sub_kernel_config = cast(KernelEnqueueingConfig,
{**kernel_enqueue_configs[0]})
            sub_kernel_config['cluster_role'] = 'sub'
            sub_kernel_config['cluster_idx'] = i + 1
            kernel_enqueue_configs.append(sub_kernel_config)
        ...
    ...
```

# Manager registry

```
for kernel in kernel_enqueue_configs:
    kernel_id: KernelId
    if kernel['cluster_role'] == DEFAULT_ROLE:
        kernel_id = cast(KernelId, session_id)
    else:
        kernel_id = KernelId(uuid.uuid4())
    creation_config = kernel['creation_config']
    image_ref = kernel['image_ref']
    ...
    image_info = await self.shared_config.inspect_image(image_ref)
    image_min_slots, image_max_slots = \
        await self.shared_config.get_image_slot_ranges(image_ref)
    # Parse service ports to check for port errors
    parse_service_ports(image_info['labels'].get('ai.backend.service-ports', ''),
BackendError)
    ...
    # Check if: requested <= image-maximum
    if not (requested_slots <= image_max_slots):
        raise InvalidAPIParameters(...)
    environ = kernel_enqueue_configs[0]['creation_config'].get('environ') or {}
```

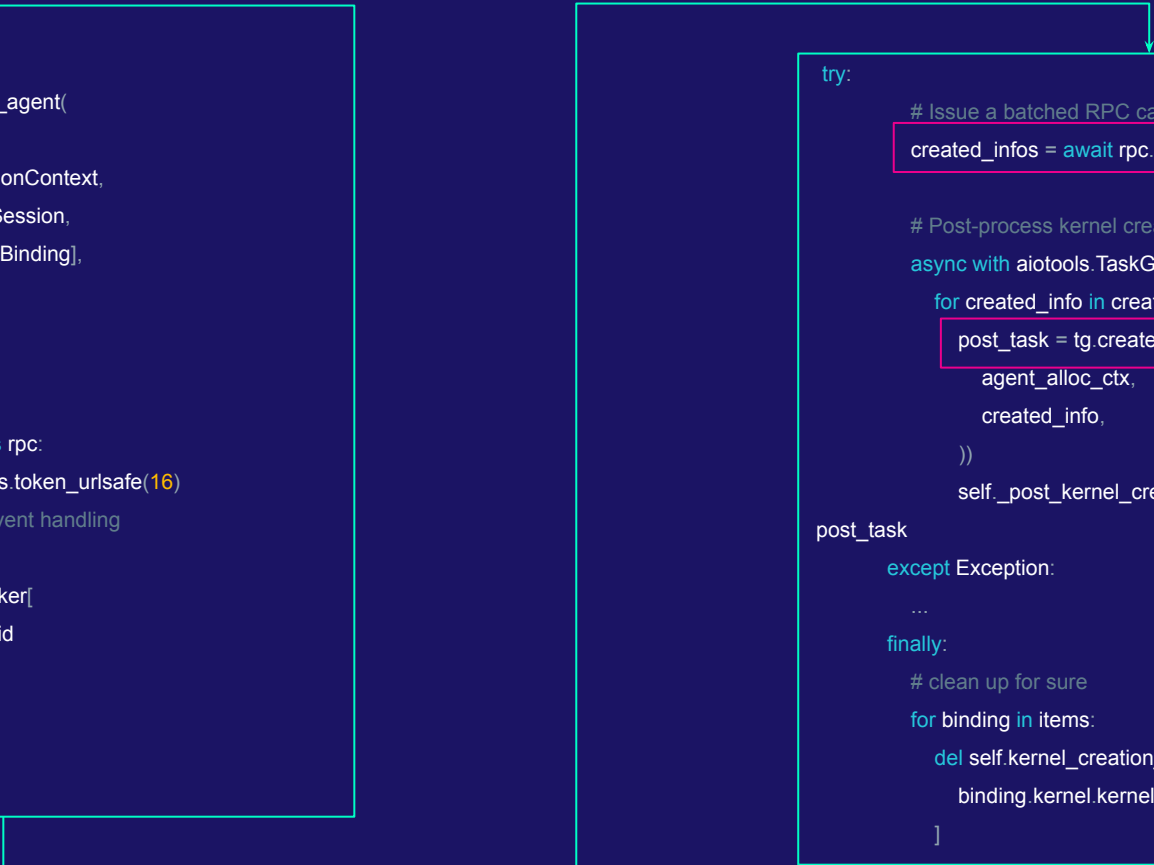


```
async def _enqueue() -> None:
    nonlocal ids
    async with self.db.begin() as conn:
        query = kernels.insert().values({
            'id': kernel_id,
            'status': KernelStatus.PENDING,
            'session_creation_id': session_creation_id,
            'session_id': session_id,
            'session_name': session_name,
            ...
            'user_uuid': user_uuid,
            'access_key': access_key,
            'image': image_ref.canonical,
            'registry': image_ref.registry,
            'tag': session_tag,
            ...
        })
        await conn.execute(query)
        ids.append(kernel_id)

    await execute_with_retry(_enqueue)
```

# Manager API session

```
async def _create_kernels_in_one_agent(
    self,
    agent_alloc_ctx: AgentAllocationContext,
    scheduled_session: PendingSession,
    items: Sequence[KernelAgentBinding],
    image_info,
    cluster_info,
) -> None:
    ...
    async with RPCContext(...) as rpc:
        kernel_creation_id = secrets.token_urlsafe(16)
        # Prepare kernel_started event handling
        for binding in items:
            self.kernel_creation_tracker[
                binding.kernel.kernel_id
            ] = loop.create_future()
```




```
try:
    # Issue a batched RPC call to create kernels on this agent
    created_infos = await rpc.call.create_kernels(...)

    # Post-process kernel creation
    async with aiotools.TaskGroup() as tg:
        for created_info in created_infos:
            post_task = tg.create_task(self._post_create_kernel(
                agent_alloc_ctx,
                created_info,
            ))
            self._post_kernel_creation_tasks[created_info['id']] =
post_task
except Exception:
    ...
finally:
    # clean up for sure
    for binding in items:
        del self.kernel_creation_tracker[
            binding.kernel.kernel_id
        ]
```



# Manager API session

```
async def _post_create_kernel(
    self,
    agent_alloc_ctx: AgentAllocationContext,
    created_info,
):
    # Wait until the kernel_started event.
    kernel_id = KernelId(uuid.UUID(created_info["id"]))
    start_event = self.kernel_creation_tracker[kernel_id]
    try:
        await start_event
    except asyncio.CancelledError:
        log.warning("post_create_kernel(k:{}) cancelled", kernel_id)
        return
    # Record kernel access information
```



```
async def _finalize_running() -> None:
    async with self.db.begin() as conn:
        agent_host = URL(agent_alloc_ctx.agent_addr).host
        kernel_host = created_info.get('kernel_host', agent_host)
        service_ports = created_info.get('service_ports', [])
        # NOTE: created_info contains resource_spec
        query = (
            kernels.update()
            .values({
                'scaling_group': agent_alloc_ctx.scaling_group,
                'status': KernelStatus.RUNNING,
                'container_id': created_info['container_id'],
                ...
            })
            .where(kernels.c.id == created_info['id']))
        await conn.execute(query)

    await execute_with_retry(_finalize_running)
```

Manager



session_name	status	kernel_id	...
psdk-1a..	PENDING	a8411..	...

Agent

Manager

Agent



04

Agent

# Agent agent

```
async def create_kernel(
    self,
    creation_id: str,
    session_id: SessionId,
    kernel_id: KernelId,
    kernel_config: KernelCreationConfig,
    cluster_info: ClusterInfo,
    *,
    restarting: bool = False,
) -> KernelCreationResult:
    """
    Create a new kernel.
    """

    ...

    # Finally we are done.
    await self.produce_event(
        KernelStartedEvent(kernel_id, creation_id)
    )
```

```
if kernel_config['session_type'] == 'batch' and
kernel_config['cluster_role'] == 'main':
    self._ongoing_exec_batch_tasks.add(
        asyncio.create_task(
            self.execute_batch(kernel_id,
                                kernel_config['startup_command'] or "")
        )
    )
return {
    'id': KernelId(kernel_id),
    'kernel_host': str(kernel_obj['kernel_host']),
    'repl_in_port': kernel_obj['repl_in_port'],
    'repl_out_port': kernel_obj['repl_out_port'],
    'stdin_port': kernel_obj['stdin_port'], # legacy
    'stdout_port': kernel_obj['stdout_port'], # legacy
    'service_ports': service_ports,
    'container_id': kernel_obj['container_id'],
    'resource_spec': resource_spec.to_json_serializable_dict(),
    'attached_devices': attached_devices,
}
```

# Agent docker agent

```
async with closing_async(Docker()) as docker:
    try:
        container = await docker.containers.create(
            config=container_config, name=kernel_name)
        cid = container._id

        resource_spec.container_id = cid
        # Write resource.txt again to update the container id.
        with open(self.config_dir / 'resource.txt', 'w') as f:
            await loop.run_in_executor(None, resource_spec.write_to_file, f)
        async with AsyncFileWriter(
            target_filename=self.config_dir / 'resource.txt',
            access_mode='a',
        ) as writer:
            for dev_name, device_alloc in resource_spec.allocations.items():
                computer_ctx = self.computers[dev_name]
                kvpairs = \
                    await
                computer_ctx.instance.generate_resource_data(device_alloc)
                for k, v in kvpairs.items():
                    await writer.write(f'{k}={v}\n')

        await container.start()
```

```
kernel_obj = await DockerKernel.new(
    self.kernel_id,
    self.image_ref,
    self.kspec_version,
    agent_config=self.local_config,
    service_ports=service_ports,
    resource_spec=resource_spec,
    data={
        'container_id': container._id,
        'kernel_host': advertised_kernel_host or container_bind_host,
        'repl_in_port': repl_in_port,
        'repl_out_port': repl_out_port,
        'stdin_port': stdin_port, # legacy
        'stdout_port': stdout_port, # legacy
        'host_ports': host_ports,
        'domain_socket_proxies': self.domain_socket_proxies,
        'block_service_ports': self.internal_data.get('block_service_ports', False),
    })
return kernel_obj
```

# Agent docker agent

```
class DockerAgent(AbstractAgent[DockerKernel,  
DockerKernelCreationContext]):
```

```
    monitor_docker_task: asyncio.Task  
    agent_sockpath: Path  
    agent_sock_task: asyncio.Task  
    scan_images_timer: asyncio.Task
```

```
    def __init__(  
        self,  
        etcd: AsyncEtcd,  
        local_config: Mapping[str, Any],  
        *,  
        stats_monitor: StatsPluginContext,  
        error_monitor: ErrorPluginContext,  
        skip_initial_scan: bool = False,  
    ) -> None:  
        ...
```

## Agent Agent

```
@abstractmethod  
    async def pull_image(self, image_ref: ImageRef, registry_conf:  
ImageRegistry) -> None:  
        """  
        Pull the given image from the given registry.  
        """
```

```
    async def pull_image(self, image_ref: ImageRef, registry_conf:  
ImageRegistry) -> None:  
        auth_config = None  
        reg_user = registry_conf.get('username')  
        reg_passwd = registry_conf.get('password')  
        if reg_user and reg_passwd:  
            encoded_creds = base64.b64encode(  
                f'{reg_user}:{reg_passwd}'.encode('utf-8')) \  
                .decode('ascii')  
            auth_config = {  
                'auth': encoded_creds,  
            }  
        log.info('pulling image {} from registry', image_ref.canonical)  
        async with closing_async(Docker()) as docker:  
            await docker.images.pull(  
                image_ref.canonical,  
                auth=auth_config)
```

Manager



start event

Agent

AGENT DOCKER

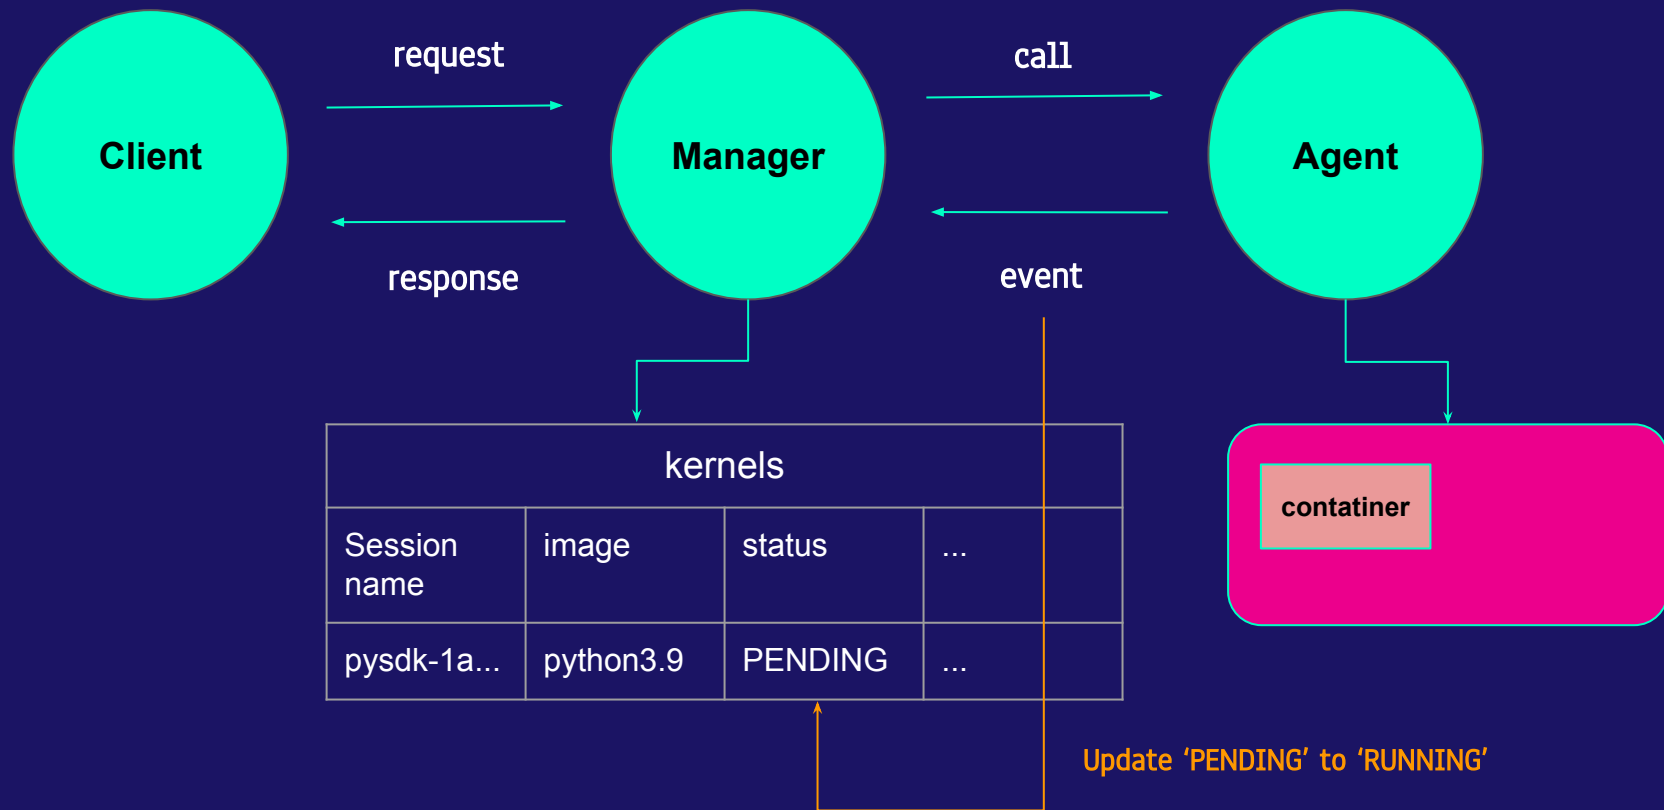
container



Agent

AGENT DOCKER

# Summary



# THANKS!

CREDITS: This presentation template was created by Slidesgo, including icons by Flaticon, and infographics & images by Freepik.